

Automation Developer



INNOVATION
SCOPE
LEARN. EMPOWER. INNOVATE

WORKFORCE DEVELOPMENT



Content Usage Parameters

Content refers to material including instructor guides, student guides, lab guides, lab or hands-on activities, computer programs, etc. designed for use in a training program



1

Content is subject to
copyright protection

2

Content may only be
leveraged by students
enrolled in the training
program

3

Students agree not to
reproduce, make
derivative works of,
distribute, publicly perform
and publicly display in any
form or medium outside of
the training program

4

Content is intended as
reference material only to
supplement the instructor-
led training

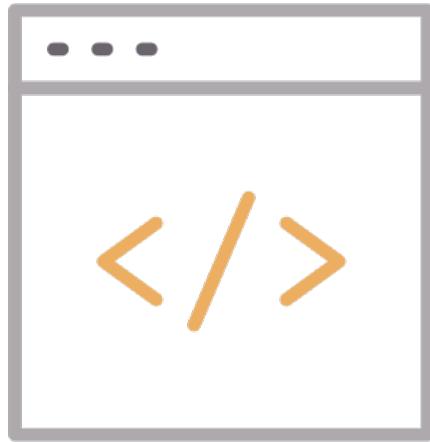
Lab page

<https://jruels.github.io/automation-dev/>



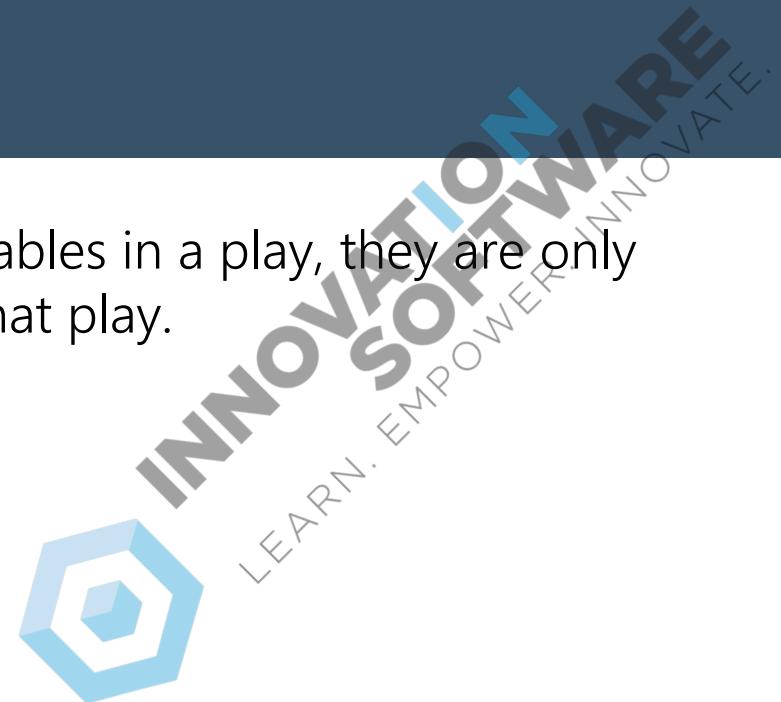
Defining Variables - Play

You can define different variables for each individual host, or set shared variables for a group of hosts in your inventory. You can also define variables in a play.



```
- hosts: webservers  
  vars:  
    http_port: 80
```

NOTE: When you define variables in a play, they are only visible to tasks executed in that play.



Defining Variables – External File



You can define variables in reusable variables files and/or in reusable roles. When you define variables in reusable variable files, the sensitive variables are separated from playbooks.

This separation enables storing playbooks in source control and even share them without exposing passwords or other sensitive data.



Defining Variables – External File

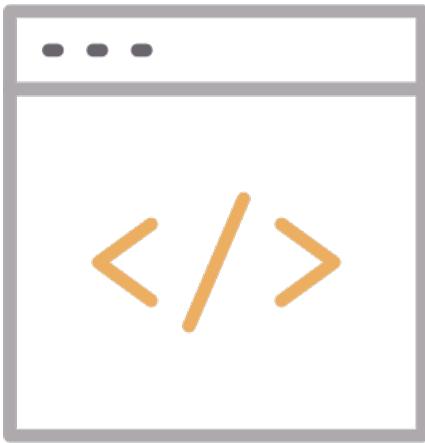


This example shows how you can include variables defined in an external file:

```
- hosts: all
  remote_user: root
  vars:
    favcolor: blue
  vars_files:
    - /var/external_vars.yml
```



Defining Variables – External File



The contents of each variables file is a simple YAML dictionary.

For example:

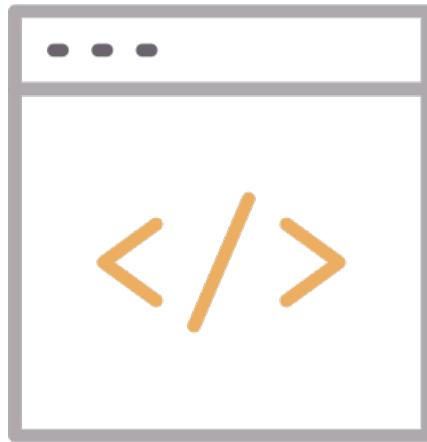
```
---  
var1: myfavorite  
var2: mysecondfavorite
```



Defining Variables – Command Line

You can define variables when you run your playbook by passing variables at the command line.

--extra-vars (-e)

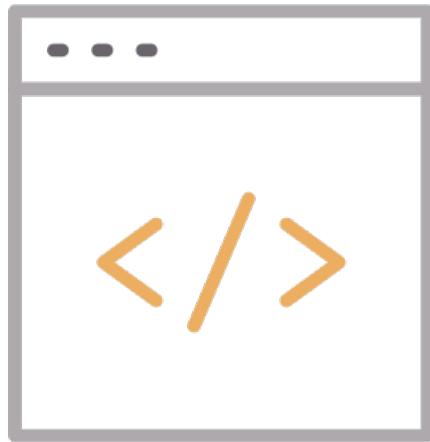


```
ansible-playbook playbook.yml --extra-vars "red"
```



Defining Variables – Command Line

If you have a lot of special characters, use a JSON or YAML file containing the variable definitions.



```
ansible-playbook release.yml -extra-vars  
"@some_var_file.json"
```



Variable Precedence



Ansible does apply variable precedence, and you might have a use for it. Here is the order of precedence from least to greatest (the last listed variables override all other variables):

1. Command line values (for example, -u my_user, these are not variables)
2. role defaults (defined in role/defaults/main.yml)
3. Inventory file or script group vars
4. Inventory group_vars/all
5. Playbook group_vars/all
6. inventory group_vars/*
7. playbook group_vars/*
8. inventory file or script host vars
9. inventory host_vars/*
10. playbook host_vars/*

Variable Precedence



11. host facts / cached set_facts
12. play vars
13. play vars_prompt
14. play vars_files
15. role vars (defined in role/vars/main.yml)
16. block vars (only for tasks in block)
17. task vars (only for the task)
18. include_vars
19. set_facts /registered vars
20. role (and include_role params)
21. include params
22. extra vars (for example –e "user=my_user") – always wins precedence.



INNOVATION SOFTWARE
LEARN. EMPOWER. INNOVATE.

Working With Variables - List



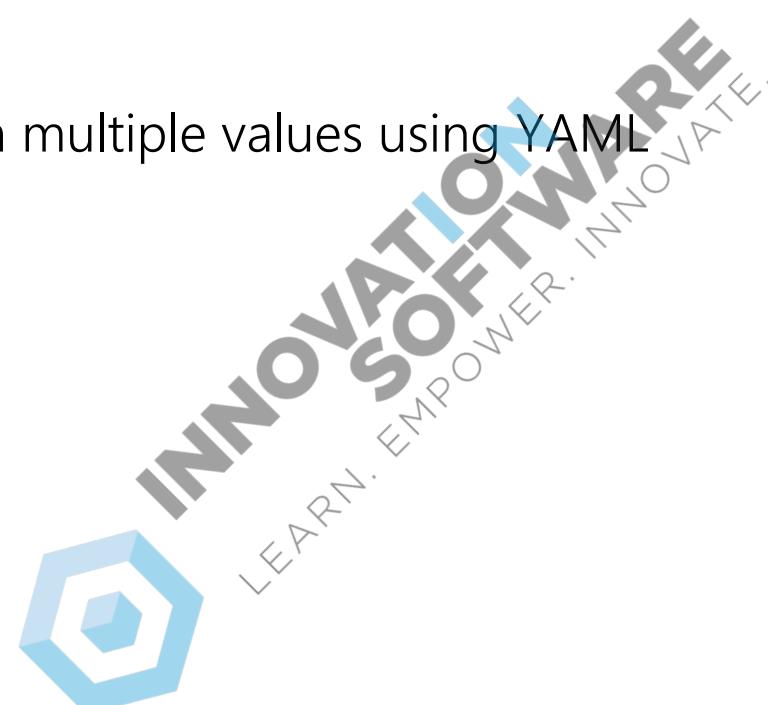
A list variable combines a variable name with multiple values. The multiple values can be stored as an itemized list or in square brackets [], separated with commas.

Defining variables as lists:

You can define variables with multiple values using YAML lists. For example:

Region:

- northeast
- southeast
- midwest



Working With Variables - List

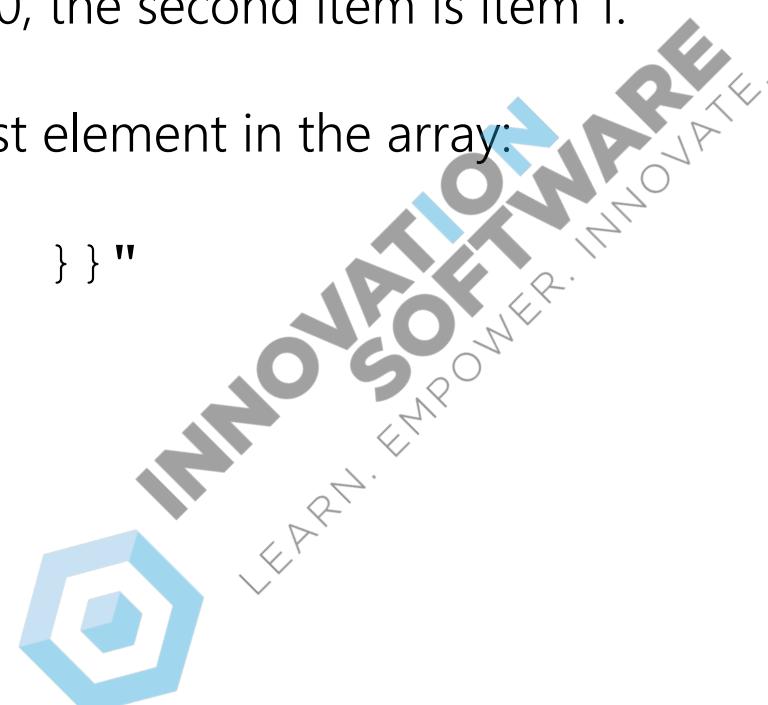


Referencing list variables:

When you use variables defined as a list (also called an array), you can use individual, specific fields from that list. The first item in a list is item 0, the second item is item 1.

For example, to select the first element in the array:

```
region: "{{ region[0] }}"
```



POP QUIZ: DISCUSSION

What does this expression return?

```
region: "{{ region[0] }}"
```



POP QUIZ: DISCUSSION

What does this expression return?

```
region: "{{ region[0] }}"
```

- A: northeast
- B: southeast
- C: midwest



POP QUIZ: DISCUSSION

What does this expression return?

```
region: "{{ region[0] }}"
```

- A: northeast
- B: southeast
- C: midwest



Working With Variables - Dictionary



A dictionary stores the data in key-value pairs. Usually, dictionaries are used to store related data, such as the information contained in an ID or a user profile.

You can define more complex variables using YAML dictionaries. A YAML dictionary maps keys to values. For example:

```
region:  
  field1: one  
  field2: two
```



Working With Variables - Dictionary



When you use variables defined as a key:value dictionary (also called a hash), you can use individual, specific fields from that dictionary using either bracket notation or dot notation:

```
foo['field']  
foo.field
```

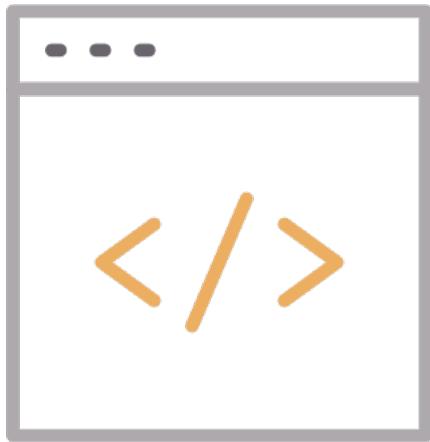
Dot notation can cause problems because some keys collide with attributes and methods of python dictionaries.

PROTIP: Use bracket notation in most cases.



Registering Variables

You can create variables from the output of an Ansible task with the task keyword `register`. You can use registered variables in any later tasks in your play.



```
- hosts: web_servers
  tasks:
    - name: Register shell command output as variable
      shell: /usr/bin/foo
      register: foo_result
      ignore_errors: true

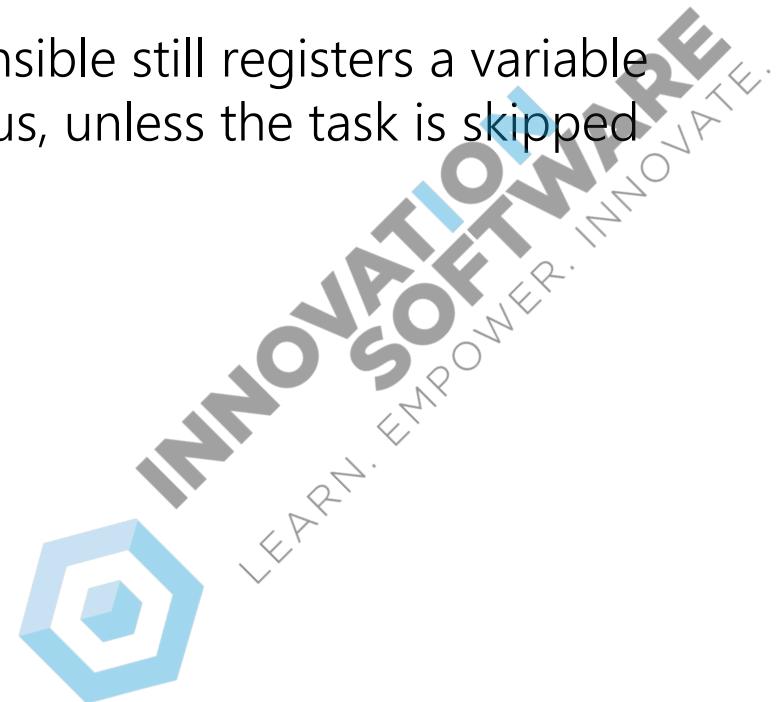
    - name: Run shell command using output from previous task
      shell: /usr/bin/bar
      when: foo_result.rc == 5
```

Registering Variables



Registered variables are stored in memory. You cannot cache registered variables for use in future plays. Registered variables are only valid on the host for the rest of the current playbook run.

If a task fails or is skipped, Ansible still registers a variable with a failure or skipped status, unless the task is skipped based on tags.

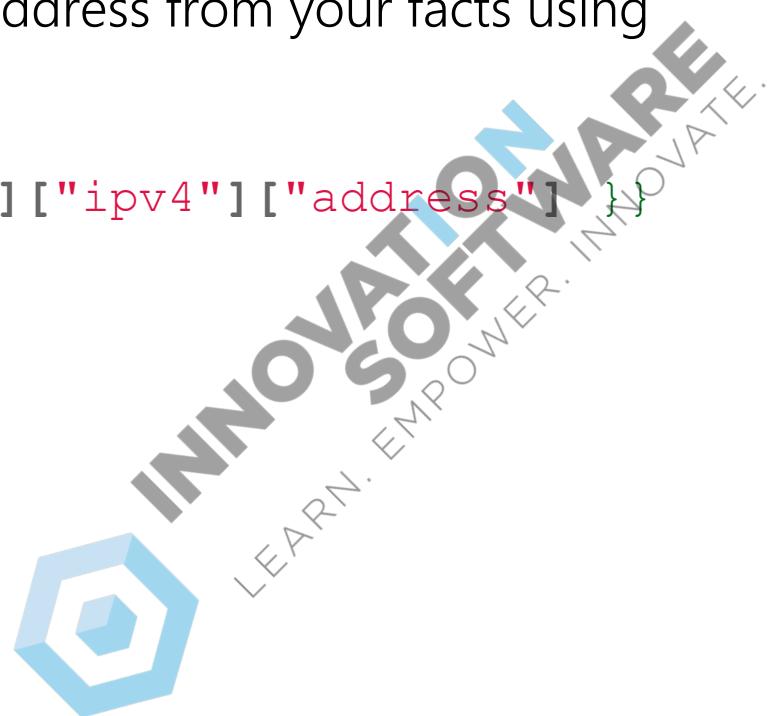


Querying Nested Data



Many registered variables (and facts) are nested YAML or JSON data structures. You cannot access values from these nested data structures with the simple {{ foo }} syntax. You must use either bracket notation or dot notation. For example, to reference an IP address from your facts using the bracket notation:

```
{ { ansible_facts["eth0"]["ipv4"]["address"] } }
```



Optional Variables

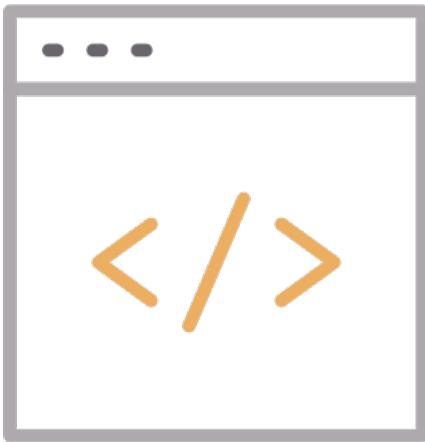


By default, Ansible requires values for all variables in a templated expression. However, you can make specific variables optional.

For example, you might want to use a system default for some items and control the value for others. To make a variable optional, set the default value to the special variable `omit`:



Optional Variables



In this example, the default mode for the files `/tmp/foo` and `/tmp/bar` is determined by the umask of the system.

Ansible does not send a value for `mode`. Only the third file, `/tmp/baz`, receives the `mode=0444` option.

```
- name: Touch files with optional mode
  ansible.builtin.file:
    dest: "{{ item.path }}"
    state: touch
    mode: "{{ item.mode | default('omit') }}"
  loop:
    - path: /tmp/foo
    - path: /tmp/bar
    - path: /tmp/baz
      mode: "0444"
```

Handlers



Sometimes you want a task to run only when a change is made on a machine. For example, you may want to restart a service if a task updates the configuration of that service, but not if the configuration is unchanged. Ansible uses handlers to address this use case. Handlers are tasks that only run when notified.



POP QUIZ: DISCUSSION

How have you used handlers?



POP QUIZ: DISCUSSION

When should handlers be used?

- When a configuration file is updated, and the service needs to be restarted.
- When a new release is rolled out



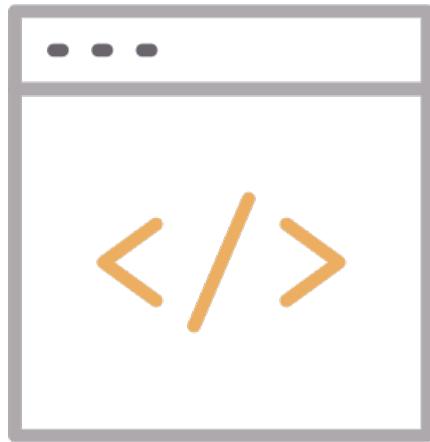
Handlers

```
---
```

```
- name: Verify Apache installation
  hosts: web
  vars:
    http_port: 80
    max_clients: 200
  tasks::
    - name: Install Apache
    - yum:
    - name: httpd
...
- name: Apache config
  template:
    src: httpd.j2
    dest: /etc/httpd.conf
  notify:
    - Restart Apache
  handlers:
    - name: Restart Apache
      service:
        name: httpd
        state: restarted
```



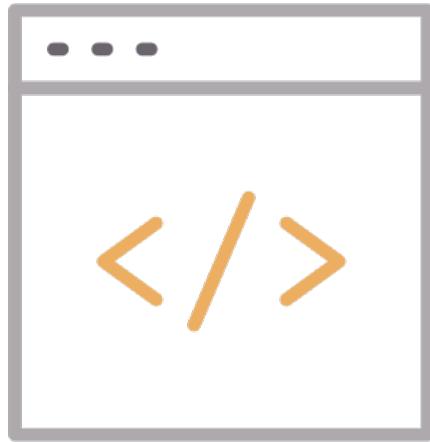
Notify Handlers



Tasks can instruct one or more handlers to execute using the `notify` keyword. The `notify` keyword can be applied to a task and accepts a list of handler names that are notified on a task change.

```
tasks:  
  name: Template file  
  template:  
    src: template.j2  
    dest: /etc/foo.conf  
  notify:  
    - Restart apache  
    - Restart memcached
```

Notify Handlers



Tasks can instruct one or more handlers to execute using the `notify` keyword. The `notify` keyword can be applied to a task and accepts a list of handler names that are notified on a task change.

```
tasks:  
  notify:  
    - Restart apache  
    - Restart memcached  
handlers:  
  - name: Restart memcached  
    service:  
      name: memcached  
      state: restarted  
  - name: Restart Apache  
    service:  
      name: apache  
      state: restarted
```

Specifying Handlers



Handlers must be named in order for tasks to be able to notify them using the `notify` keyword.

Alternately, handlers can utilize the `listen` keyword. Using this handler keyword, handlers can listen on topics that can group multiple handlers as follows:



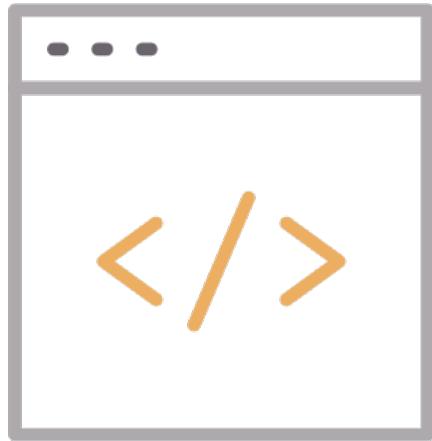
Specifying Handlers

```
tasks:::  
  - name: Restart everything  
    command: echo "this task will restart the web services"  
    notify: "restart web services"  
  
handlers:  
  - name: Restart memcached  
    service:  
      name: memcached  
      state: restarted  
      listen: "restart web services"  
  
  - name: Restart Apache  
    service:  
      name: httpd  
      state: restarted  
      listen: "restart web services"
```



INNOVATION
SOFTWARE
LEARN. EMPOWER. INNOVATE.

Blocks



Example with Ansible Blocks

```
---
```

```
- name: Ansible Blocks
  hosts: server1
  gather_facts: false

  tasks:
    - block:
        - name: List usr directory
          command: "ls -l /usr/"

        - name: List root directory
          command: "ls -l /root"
          become: yes

        - name: List home directory
          command: "ls -l ~/"
```

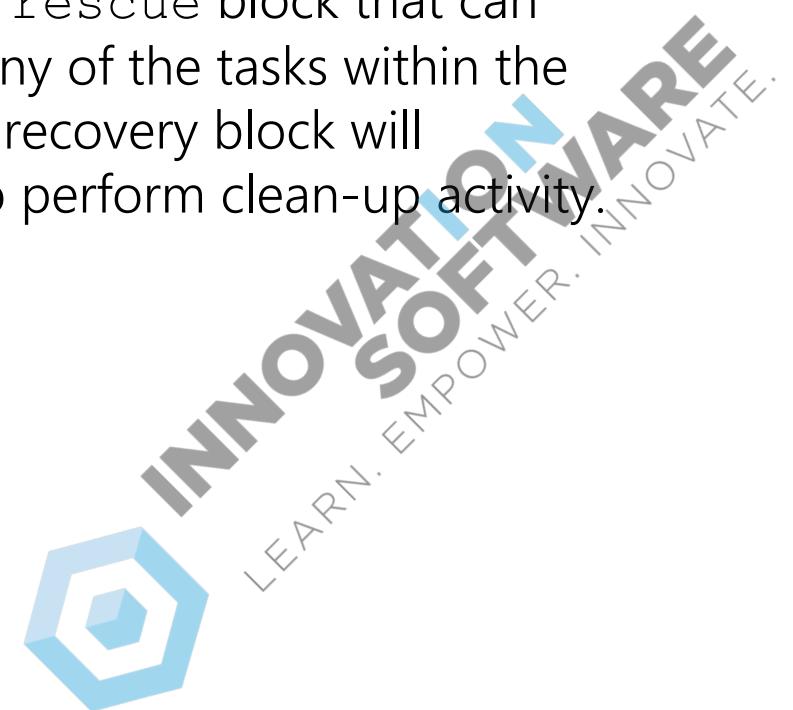
Recovery



An additional benefit of using ansible blocks is to perform recovery operations.

If any of the tasks within a block fail, the playbook will exit.

With blocks, we can assign a `rescue` block that can contain a bunch of tasks. If any of the tasks within the block fail, the tasks from the recovery block will automatically be executed to perform clean-up activity.



Rescue Variables



Ansible provides a couple of variables for tasks in the rescue portion of a block:

- `ansible_failed_task`

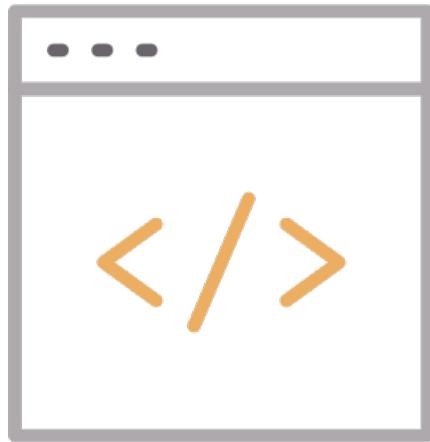
The task that returned 'failed' and triggered the rescue. For example, to get the name use `ansible_failed_task.name`

- `ansible_failed_result`

The captured return result of the failed task that triggered the rescue. The same as registering the variable.



Rescue Block



Example with Rescue block

```
---
```

```
- name: Ansible Blocks
  hosts: server1
  gather_facts: false

  tasks:
    - block:
        - name: List home directory
          command: "ls -l ~/"

        - name: Failing intentionally
          command: "ls -l /tmp/not-home"

    rescue:
      - name: Rescue block (perform recovery)
        debug:
          msg: "Something broke! Cleaning up..."
```

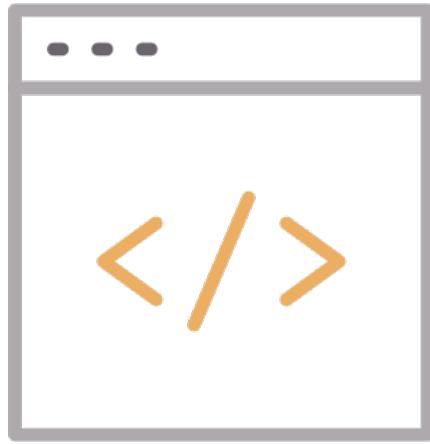
Always Block



An always block will be called independent of the task execution status. It can be used to give a summary or perform additional tasks whether the block tasks fail or not.



Always Block



Example with Always block

```
---
```

```
- name: Ansible Blocks
  hosts: server1
  gather_facts: false

  tasks:
    - block:
        - name: List home directory
          command: "ls -l ~/"

        - name: Failing intentionally
          command: "ls -l /tmp/not-home"

    always:
      - name: This always executes
        debug:
          msg: "Can't stop me..."
```

Block Practical Example



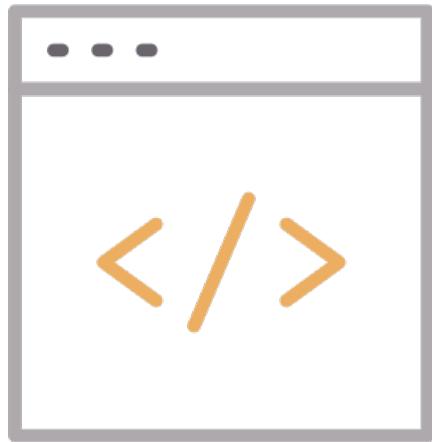
Now that we've discussed how a block can be used, let's look at practical examples.

- Install, configure, and start a service
- Apply logic to all tasks in the block
- Enable error handling



Block Practical Example

Practical example (Install, configure, and start Apache)



```
---
```

```
tasks:
```

```
  - name: Install, configure, start Apache
```

```
  block:
```

```
    - name: Install httpd and memcached
```

```
      ansible.builtin.yum:
```

```
        name:
```

```
          - httpd
```

```
          - memcached
```

```
        state: present
```

```
    - name: Apply config template
```

```
      ansible.builtin.template:
```

```
        src: templates/src.j2
```

```
        dest: /etc/template.conf
```

```
    - name: Start/enable service
```

```
      ansible.builtin.service:
```

```
        name: httpd
```

```
        state: started
```

```
        enabled: true
```

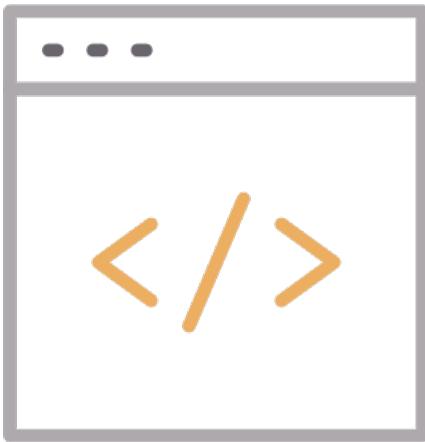
```
        when: ansible_facts['distribution'] == 'CentOS'
```

```
        become: true
```

```
        become_user: root
```

Block Practical Example

The when condition evaluated for all tasks in block.



Practical example (Install, configure, and start Apache

```
---
```

```
tasks:
```

```
  - name: Install, configure, start Apache
```

```
    block:
```

```
      - name: Install httpd and memcached
```

```
        ansible.builtin.yum:
```

```
          name:
```

```
            - httpd
```

```
            - memcached
```

```
          state: present
```

```
      - name: Apply config template
```

```
        ansible.builtin.template:
```

```
          src: templates/src.j2
```

```
          dest: /etc/template.conf
```

```
      - name: Start/enable service
```

```
        ansible.builtin.service:
```

```
          name: httpd
```

```
          state: started
```

```
          enabled: true
```

```
        when: ansible_facts['distribution'] == 'CentOS'
```

```
        become: true
```

```
        become_user: root
```

Block Rescue Task Status

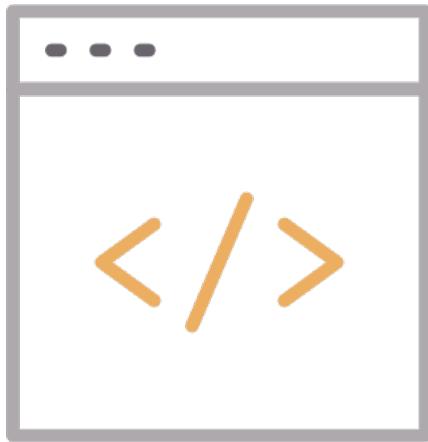


If an error occurs in the block and the rescue task succeeds, Ansible reverts the failed status of the original task for the run and continues to run the play as if the original task had succeeded.

The rescued task is considered successful. However, Ansible still reports a failure in the playbook statistics.



Block Handlers



You can use blocks with `flush_handlers` in a rescue task to ensure that all handlers run even if an error occurs:

```
---
```

```
tasks:
  - name: Attempt graceful rollback
    block:
      - name: Print a message
        ansible.builtin.debug:
          msg: 'I execute normally'
        changed_when: yes
        notify: run me even after an error

      - name: Force a failure
        ansible.builtin.command: /bin/false
    rescue:
      - name: Make sure all handlers run
        meta: flush_handlers

handlers:
  - name: Run me even after an error
    ansible.builtin.debug:
      msg: 'This handler runs even on error'
```

Lab: Ansible error handling

Jinja2 Templates



Jinja2 templates are simple template files that store variables that can change from time to time. When Playbooks are executed, these variables get replaced by actual values defined in Ansible Playbooks. This way, templating offers an efficient and flexible solution to create or alter configuration file with ease.



Jinja2 Templates



A Jinja2 template file is a text file that contains variables that get evaluated and replaced by actual values upon runtime or code execution. In a Jinja2 template file, you will find the following tags:

`{ { } }` : These double curly braces are the widely used tags in a template file and they are used for embedding variables and ultimately printing their value during code execution.

`{ % }` : These are mostly used for control statements such as loops and if-else statements.

`{ # }` : These denote comments that describe a task.



Jinja2 Templates



In most cases, Jinja2 template files are used for creating files or replacing configuration files on servers.

Apart from that, you can perform conditional statements such as `loops` and `if-else` statements and transform the data using filters and so much more.

Template files have the `.j2` extension, implying that Jinja2 templating is in use.



Jinja2 Templates

A simple Jinja2 template example.

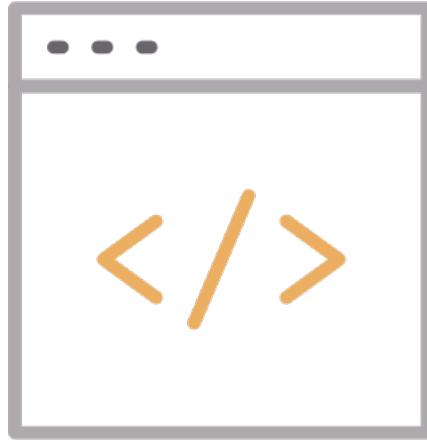
Hey guys! Apache webserver {{ version_number }} is running on {{ server }} Enjoy!



The variables are "{{ version_number }}" and "{{ server }}"



Jinja2 Templates



When the playbook is executed, the variables in the template file are replaced with declared vars.

```
---
```

```
- hosts:
```

```
  vars:
```

```
    version_number: "2.3.52"
```

```
    server: "Ubuntu"
```

```
tasks:
```

```
  - name: Jinja 2 template example
```

```
    template:
```

```
      src: my_template.j2
```

```
      dest: /home/ansible/myfile.txt
```



Lab: Ansible templates

Asynchronous Actions & Polling

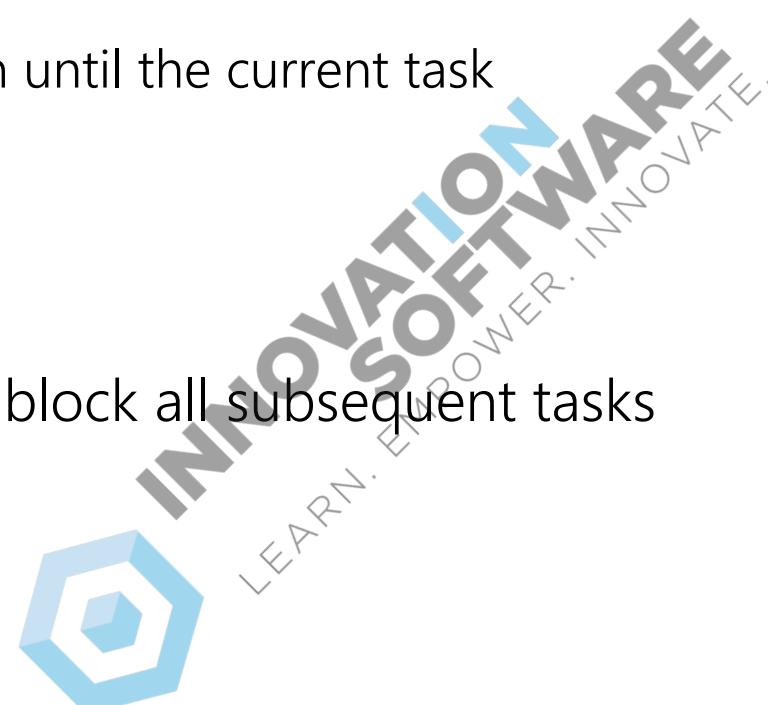


Ansible runs tasks synchronously, holding the connection to the remote node open until the action is completed. This means within a playbook; each task blocks the next task.

Subsequent tasks will not run until the current task completes.

Challenges:

- Slow
- Long running tasks block all subsequent tasks



Asynchronous Actions & Polling



Playbooks support asynchronous mode and polling, with a simplified syntax.

You can use asynchronous mode in playbooks to avoid connection timeouts or to avoid blocking subsequent tasks. The behavior of asynchronous mode in a playbook depends on the value of poll.



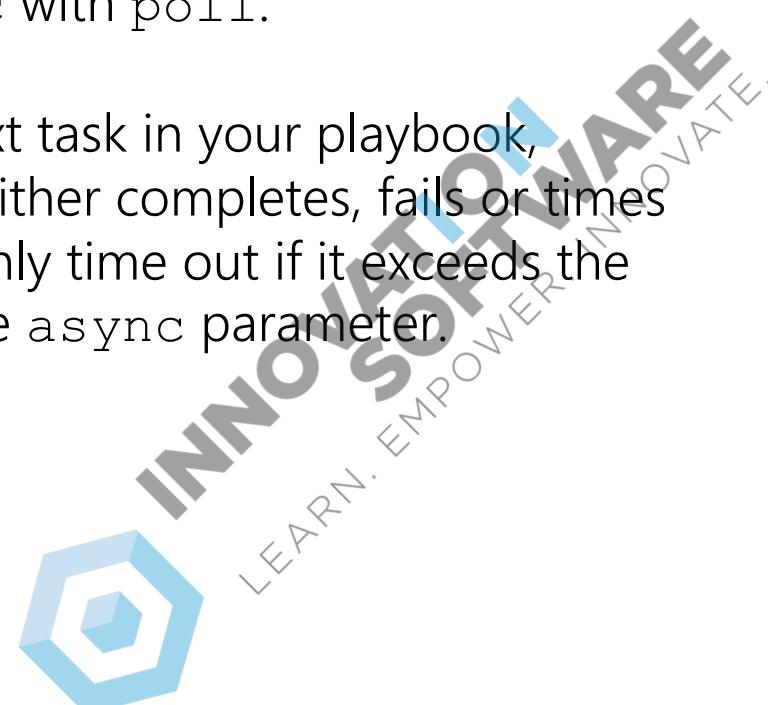
Asynchronous Actions & Polling



For long running tasks, connections to the host can timeout.

If you want to set a longer timeout limit for a certain task in your playbook, use `async` with `poll`.

Ansible will still block the next task in your playbook, waiting until the `async` task either completes, fails or times out. However, the task will only time out if it exceeds the timeout limit you set with the `async` parameter.



Asynchronous Actions & Polling

To avoid timeouts on a task, specify its maximum runtime and how frequently you would like to poll for status:

```
---
```

```
tasks:
  - name: Long running task (15 sec), wait for up to 45 sec, poll every 5 sec
    command: /bin/sleep 15
    async: 45
    poll: 5
```

- DEFAULT_POLL_INTERVAL
The default polling value is 15 seconds.

There is no default for the async time limit. If you omit the `async` keyword the tasks run synchronously.

Default async job cache file: `~/.ansible_async`



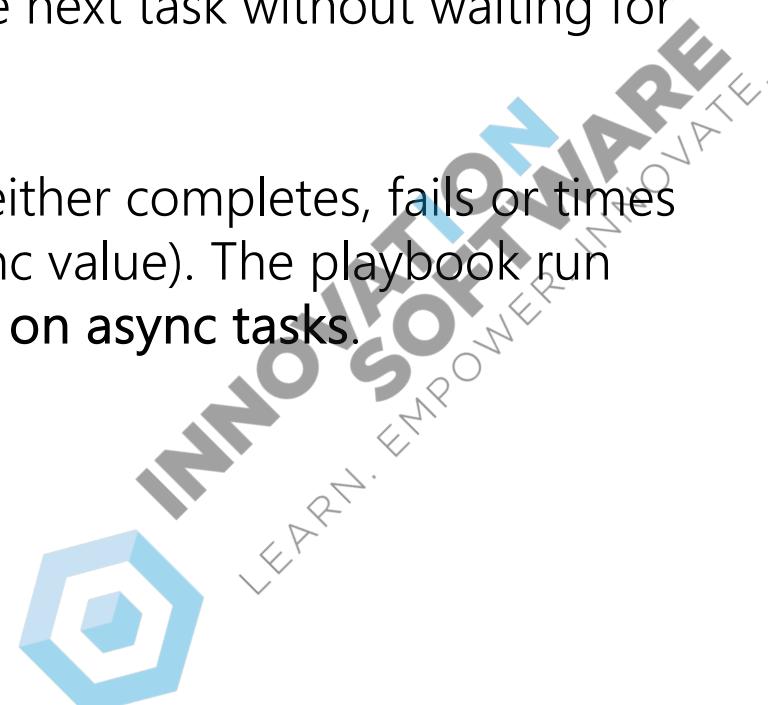
Asynchronous Actions & Polling



If you want to run multiple tasks in a playbook concurrently, use `async` with `poll` set to 0.

When you set `poll: 0`, Ansible starts the task and immediately moves on to the next task without waiting for a result.

Each `async` task runs until it either completes, fails or times out (runs longer than its `async` value). The playbook run ends **without checking back on `async` tasks**.



Asynchronous Actions & Polling

Playbook with asynchronous task:

```
---
```

```
tasks:
  - name: Long running task, allow for 45 sec, fire and forget
    command: /bin/sleep 15
    async: 45
    poll: 0
```

Be careful! Operations that require a lock (yum, apt, etc.) should not be run using `async` if you intend to run other commands later in the playbook on them.

When running with `poll: 0`, Ansible will not automatically cleanup the `async` job cache file. It will need to be cleaned up manually using the `async_status` module with mode: `cleanup`.

Asynchronous Actions & Polling

Check the status of an `async` task

```
- name: async task
  yum:
    name: docker-io
    state: present
  async: 1000
  poll: 0
  register: yum_sleeper

- name: Check status of async task
  async_status:
    jid: "{{ yum_sleeper.ansible_job_id }}"
  register: job_result
  until: job_result.finished
  retries: 100
  delay: 10
```



Check And Diff Mode



Ansible provides two modes of execution that validate tasks:

- Check mode
 - Ansible runs without making any changes on remote systems.
- Diff mode
 - Ansible provides before-and-after comparisons.



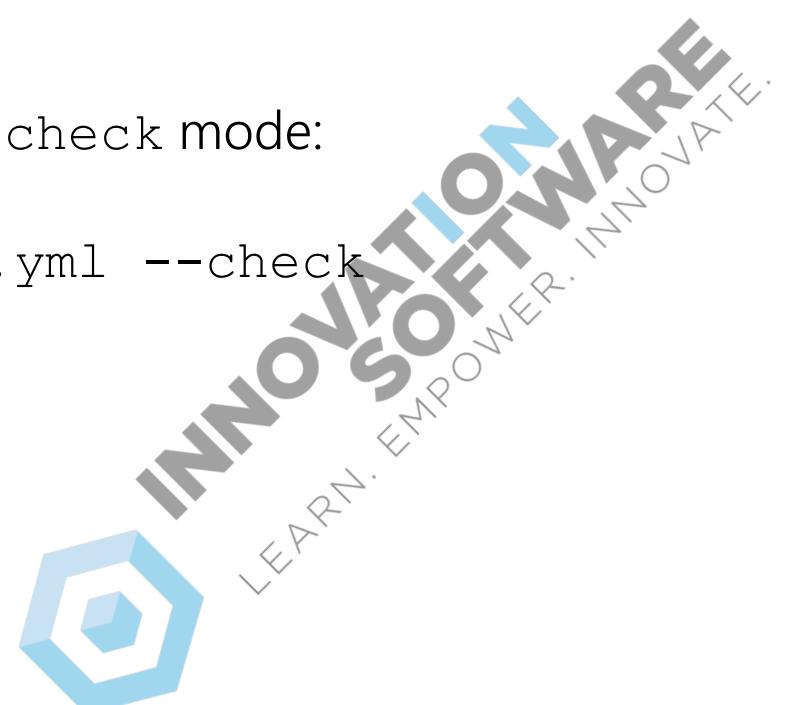
Check And Diff Mode



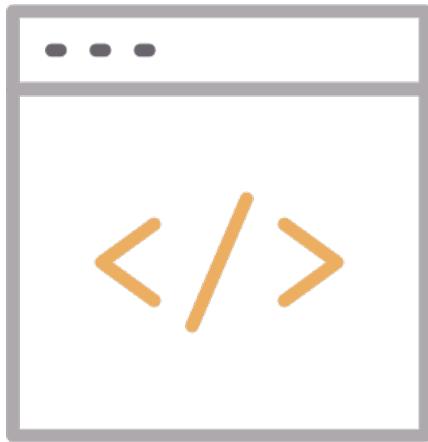
Check mode is just a simulation. It will not generate output for tasks that use conditionals based on registered variables (results of prior tasks). However, it is great for validating configuration management playbooks that run on one node at a time.

To run an entire playbook in check mode:

```
ansible-playbook foo.yml --check
```



Check And Diff Mode



It is also possible to specify that a task always or never runs in check mode regardless of command-line argument.

```
tasks:  
  - name: Always change system  
    command: /bin/change_stuff --even-in-check-mode  
    check_mode: no  
  
  - name: Never change system  
    lineinfile:  
      line: "important config"  
      dest: /path/to/config.conf  
      state : present  
    check_mode: yes  
    register: changes_to_important_config
```

Check And Diff Mode

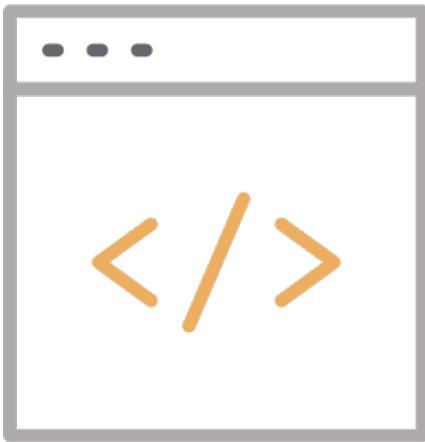


Running single tasks with `check_mode` can be useful for testing Ansible modules, either to test the module itself or to test the conditions under which it would make changes.

Combining `check_mode` and `register` provides even more detail on potential changes.



Check And Diff Mode



It is possible to skip a task or ignore errors when using `check_mode` by specifying `ansible_check_mode` boolean

```
tasks:  
  - name: Skip in check mode  
    git:  
      repo: ssh://git@github.com/me/hello-world.git  
      dest: /home/me/hello-world  
      when: not ansible_check_mode  
  
  - name: Ignore errors in check mode  
    git:  
      repo: ssh://git@github.com/me/hello-world.git  
      dest: /home/me/hello-world  
      ignore_errors: "{{ ansible_check_mode }}"
```

Check And Diff Mode



The `--diff` option for `ansible-playbook` can be used with `--check` or alone.

When you run in diff mode, any module that supports diff mode reports the changes made or, if used with `--check`, the changes that would have been made.

Diff mode is most common in modules that manipulate files (for example, the `template` module) but other modules might also show 'before and after' information (for example, the `user` module).



INNOVATION IN SOFTWARE
LEARN. EMPOWER. INNOVATE.

Why The Ansible Automation Platform?

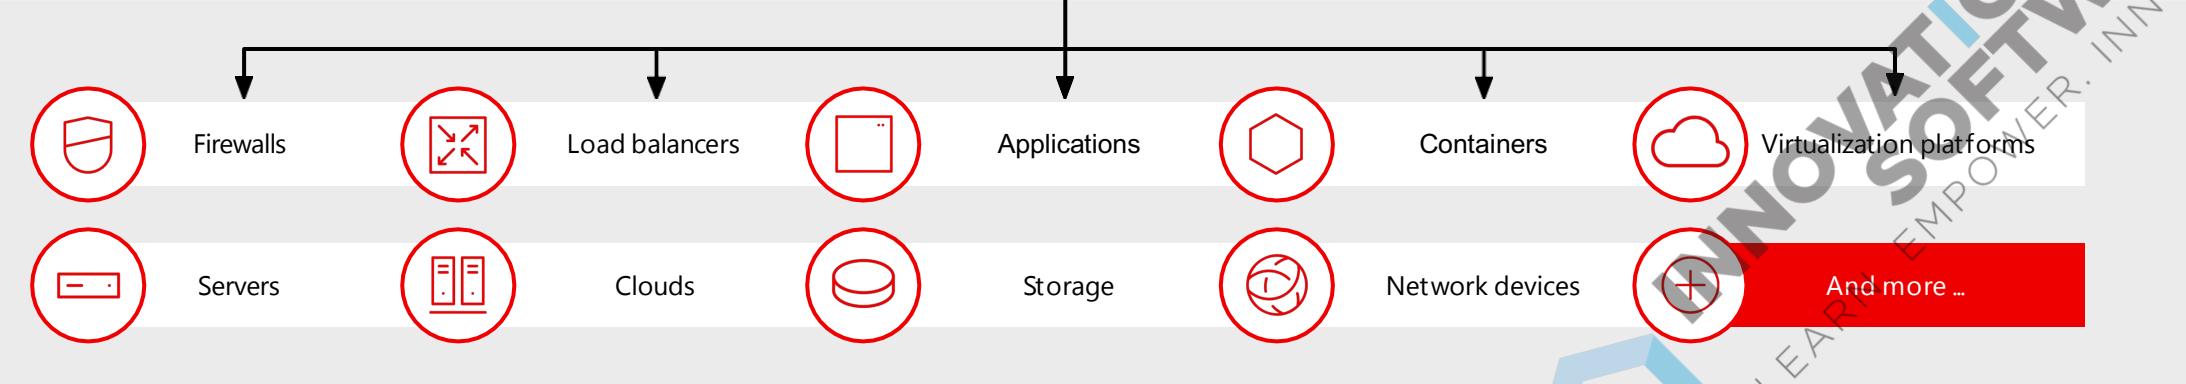
Automate the deployment and management of automation

Your entire IT footprint

Do this...

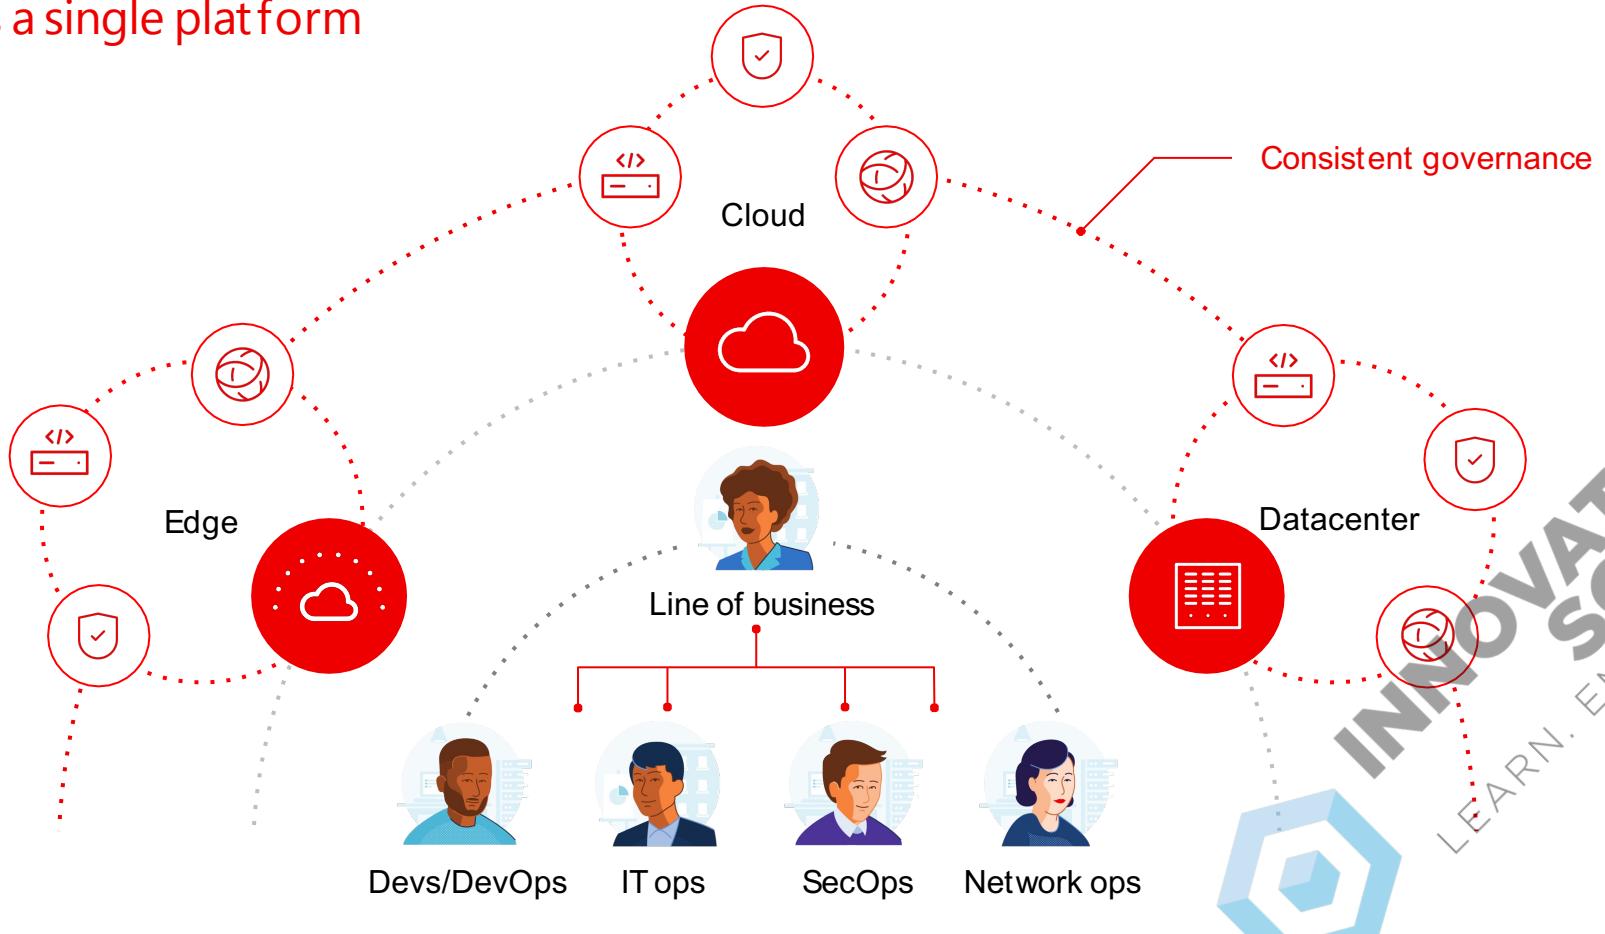
Orchestrate Manage configurations Deploy applications Provision / deprovision Deliver continuously Secure and comply

On these...



Break Down Silos

Different teams a single platform



INNOVATION SOFTWARE
LEARN. EMPOWER. INNOVATE.



Red Hat Ansible Automation Platform



Content creators



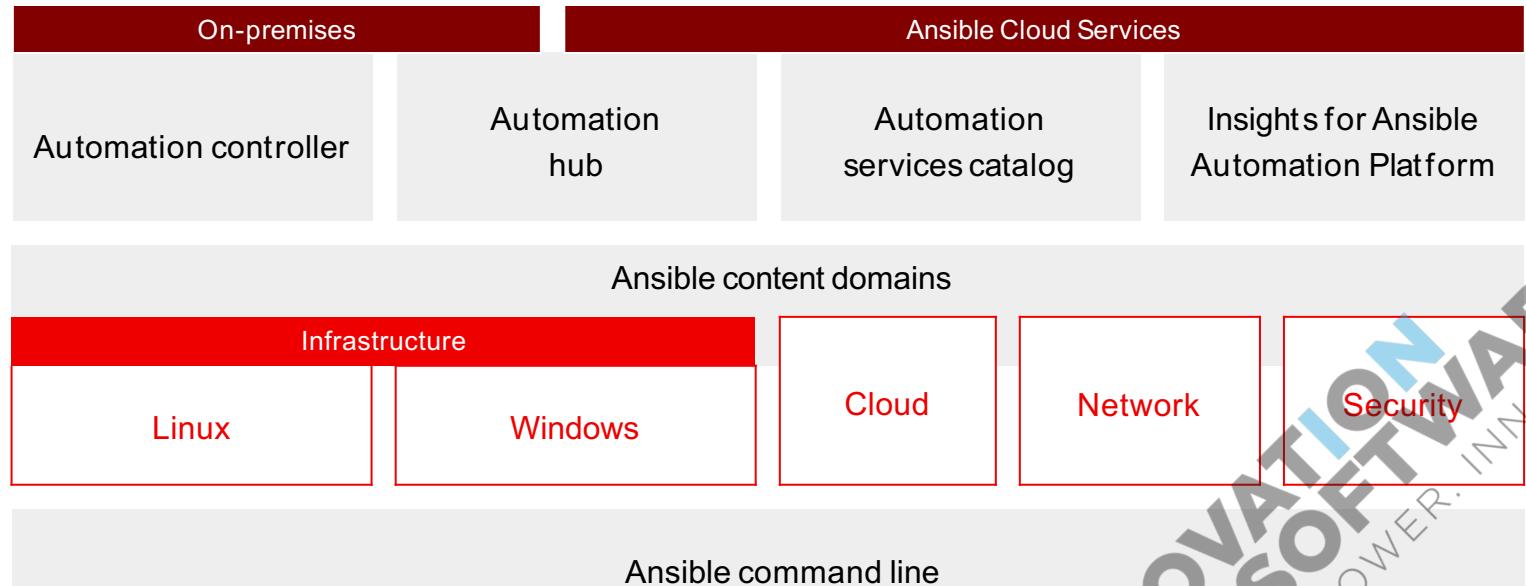
Operators



Domain experts



Users



Fueled by an
open source community



INNOVATION SOFTWARE
LEARN. EMPOWER. INNOVATE.

Automation Platform Concepts



A control node is installed with Ansible and is used to run playbooks.

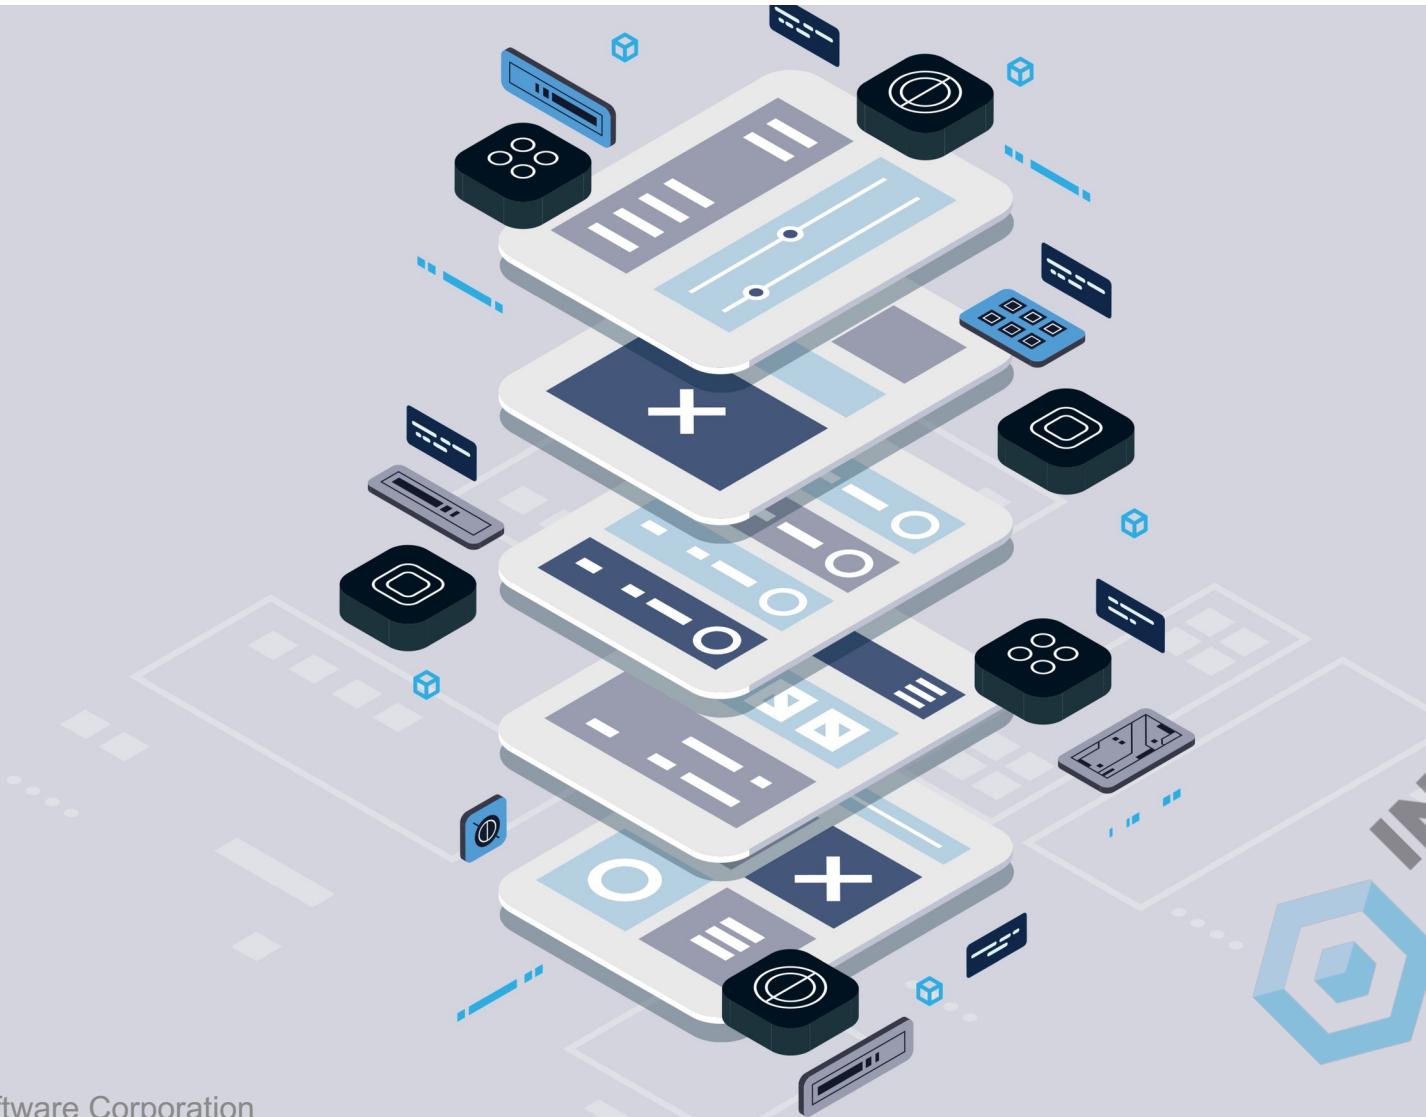
- Contains the Ansible Engine software, the playbook, and its supporting files.
- Red Hat Automation Platform is a control node that also provides a central web interface, authentication, and API for Ansible.

A managed host is a machine that is managed by Ansible automation.

- Does not have Ansible installed
- Does need to be configured to allow Ansible to connect to the host
- Must be listed in the inventory (or generated by a dynamic inventory script or plugin)



Ansible Automation Platform Infrastructure



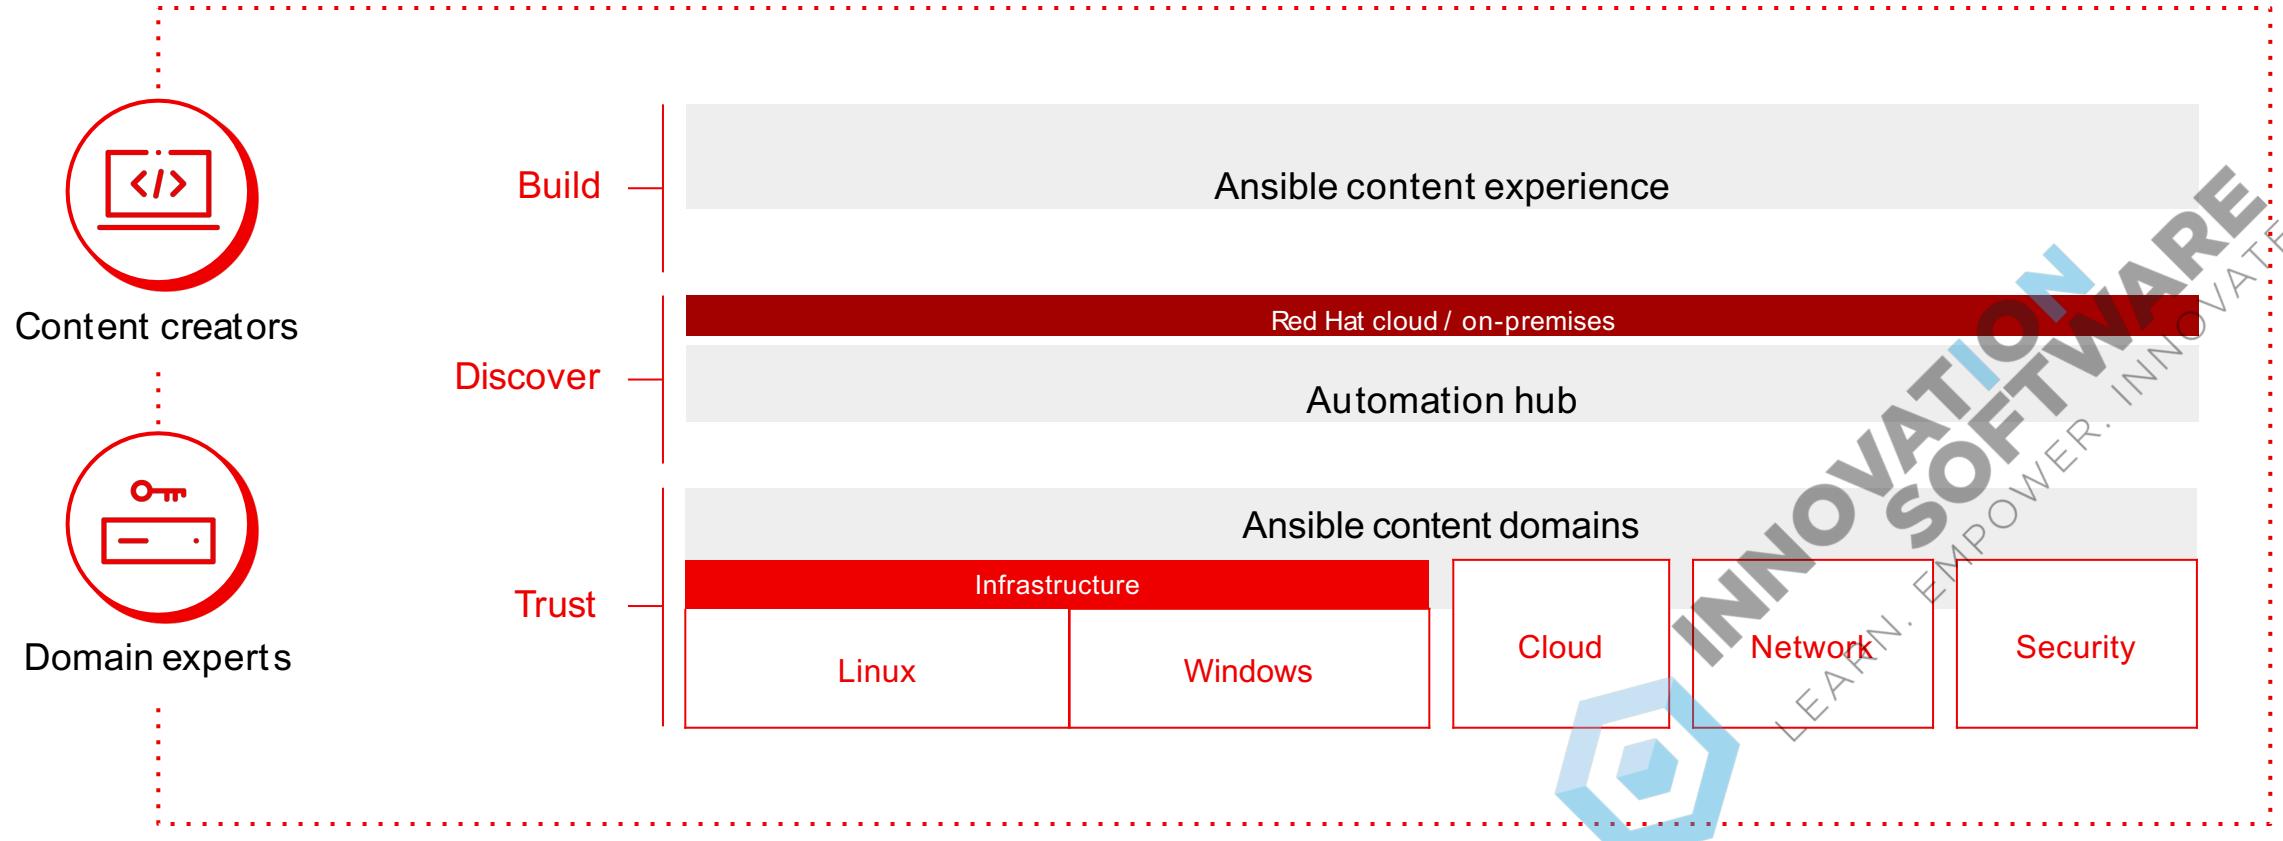
INNOVATION SOFTWARE
LEARN. EMPOWER. INNOVATE.



Create Code

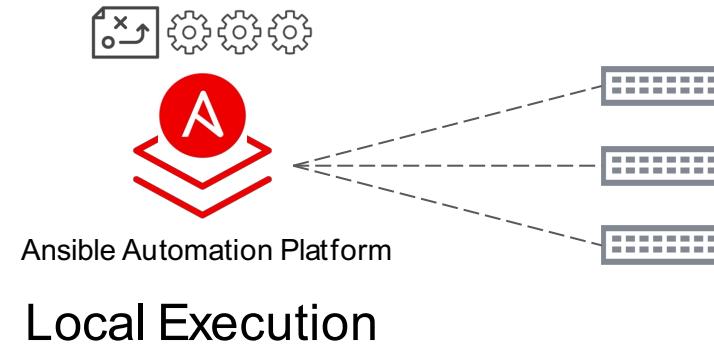
Create

The automation lifecycle



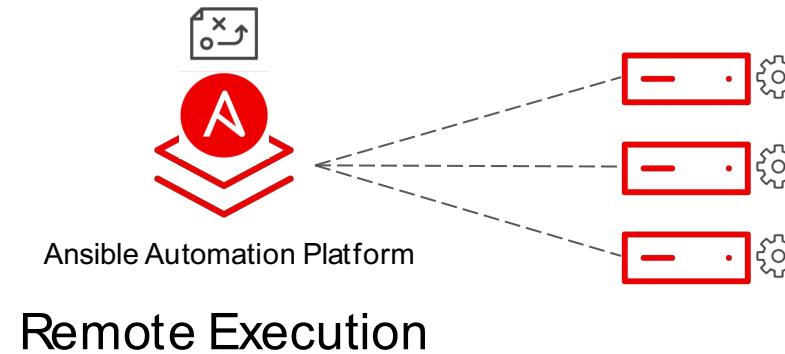
How Ansible Automation Works

Module code is executed locally on the control node



Network Devices / API Endpoints

Module code is copied to the managed node, executed, then removed



Linux / Windows Hosts

Installation Architecture Options

Red Hat Ansible Automation Platform can be implemented using one of the following architectures:

- Single Machine with Integrated Database
- Single Machine with Remote Database
- Multi Machine Cluster with Remote Database
- OpenShift Pod with Remote Database



Lab Environment

This class uses a single-node installation with an integrated database:

- Need a system installed with Red Hat Enterprise Linux (bare metal, virtual machine, or cloud instance)
 - Cloud VM
 - 2 vCPU / 4 GB RAM / at least 40GB of storage

Sign up for an evaluation of Automation Platform from Red Hat:

- <https://www.redhat.com/en/technologies/management/ansible/try-it>



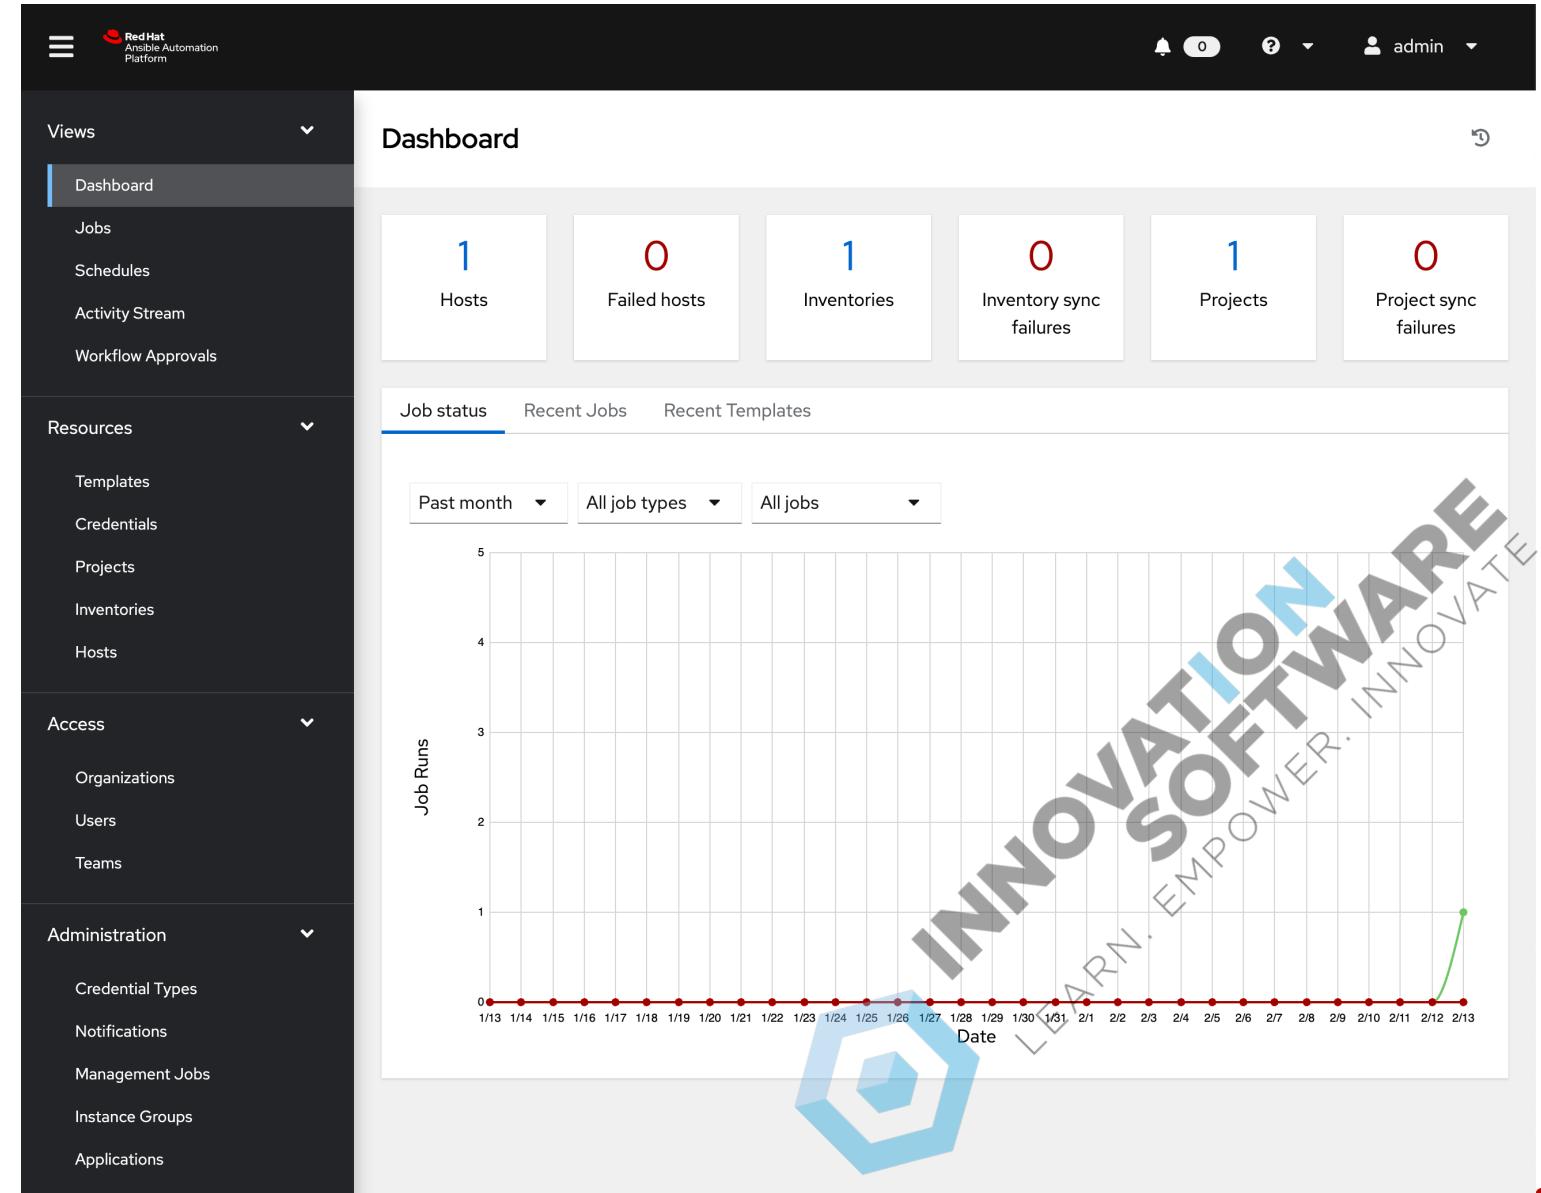
Installer

Two different installation packages are available for Automation Platform:

- Standard setup
 - <https://access.redhat.com/downloads/content/480>
 - Requires internet connectivity to download Automation Platform packages from repositories.
- Bundled installer
 - Download from same link as “Standard setup”, but choose the packages with “Bundle” in the name.
 - Includes initial RPM packages for Automation Platform
 - May be installed on systems without internet access.

Automation Platform Dashboard

- The main control center for Red Hat Ansible Tower.
- Displayed when you log in.
- Composed of four reporting sections:
 - Summary
 - Job Status
 - Recently Used
 - TemplatesRecent Job Runs



Dashboard Navigation

The dashboard contains links to common resources.

Organizations	Control what resources are visible to which users.
Teams	Groups of users that need to access the same resources
Users	User accounts
Jobs	A history of previous Ansible runs.
Templates	Prepared playbooks settings that can be "launched" to run a job.
Credentials	Authentication secrets for managed hosts and Git repositories
Projects	Sources of Ansible Playbooks (usually Git repo)
Inventories	Inventories of managed hosts
Inventory Scripts	Dynamic inventory
Notifications	Configurable for job completion or failure
Management Jobs	Special jobs used to maintain Ansible Platform.

Lab: Install Ansible Automation Platform

